

Understanding K-Nearest Neighbors (KNN)

K-Nearest Neighbors (KNN) is one of the simplest yet powerful algorithms in machine learning.  Whether you're working on classification or regression, KNN is a go-to algorithm for quick, interpretable solutions. Let's explore! 

What is KNN?

KNN is a supervised learning algorithm that predicts the output for a data point based on its K-nearest neighbors in the feature space. Think of it as finding "friends" in your data! 

How It Works:

1 Calculate Distance:

For a given data point, KNN calculates the distance to all other points in the dataset (using metrics like Euclidean, Manhattan, etc.). 

2 Find Neighbors:

Identify the K-nearest points to the input data.

3 Make Predictions:

Classification: The label is decided by majority vote among the neighbors. 

Regression: The prediction is the average of the neighbors' values. 

Choosing the Right K:

A small K may lead to overfitting. 

A large K smoothens predictions but may underfit. 

Use cross-validation to find the optimal value of K. 

Strengths:

Simple and intuitive. 

Effective for smaller datasets. 

No need for assumptions about the data distribution. 

⚠️ Limitations:

Computationally expensive for large datasets.  

Sensitive to irrelevant features and scaling. 

Choice of K and distance metric significantly impacts performance.



GitHub Code: <https://github.com/NafisAnsari786/Machine-Learning-Algorithms/blob/main/16%20KNN%20Classification/KNN%20classification%20on%20Iris%20data.ipynb>

<https://github.com/NafisAnsari786/Machine-Learning-Algorithms/blob/main/16%20KNN%20Classification/KNN%20classification%20on%20Iris%20data.ipynb>